

Kernel Methods Classification

David Park

2022-10-25

```
library(e1071)
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.3.2 --
## v ggplot2 3.3.6      v purrr  0.3.4
## v tibble  3.1.8      v dplyr  1.0.10
## v tidyr   1.2.1      v stringr 1.4.1
## v readr   2.1.3      v forcats 0.5.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

```
set.seed(1234)
```

Extract 10k observations from dataset

```
df <- read.csv("C:\\School\\CS 4375 Machine Learning\\Kernel Ensemble Methods\\Invistico_Airline.csv")
df <- sample_n(df, 10000, replace=FALSE)
df$satisfaction <- as.factor(df$satisfaction)
```

Combine columns that give rating into one

```
df$rating.mean <- as.numeric(apply(df[, 8:21], 1, sum)) / 14
df <- df[, c(1,4,24)]
str(df)
```

```
## 'data.frame': 10000 obs. of 3 variables:
## $ satisfaction: Factor w/ 2 levels "dissatisfied",...: 2 2 1 2 1 2 2 2 1 2 ...
## $ Age : int 32 43 66 69 10 36 42 46 41 26 ...
## $ rating.mean : num 3 3.14 2.79 3.57 3 ...
```

```
head(df)
```

```
## satisfaction Age rating.mean
## 1 satisfied 32 3.000000
## 2 satisfied 43 3.142857
```

```
## 3 dissatisfied 66      2.785714
## 4      satisfied 69      3.571429
## 5 dissatisfied 10      3.000000
## 6      satisfied 36      3.214286
```

Divide train/test

```
spec <- c(train=.6, test=.2, validate=.2)
i <- sample(cut(1:nrow(df), nrow(df)*cumsum(c(0,spec))), labels=names(spec))
train <- df[i=="train",]
test <- df[i=="test",]
vald <- df[i=="validate",]
```

Explore the training data statistically and graphically

```
str(train)
```

```
## 'data.frame':    6000 obs. of  3 variables:
##  $ satisfaction: Factor w/ 2 levels "dissatisfied",...: 2 2 1 2 2 2 1 2 1 1 ...
##  $ Age          : int   32 43 66 69 36 26 51 46 37 45 ...
##  $ rating.mean  : num   3 3.14 2.79 3.57 3.21 ...
```

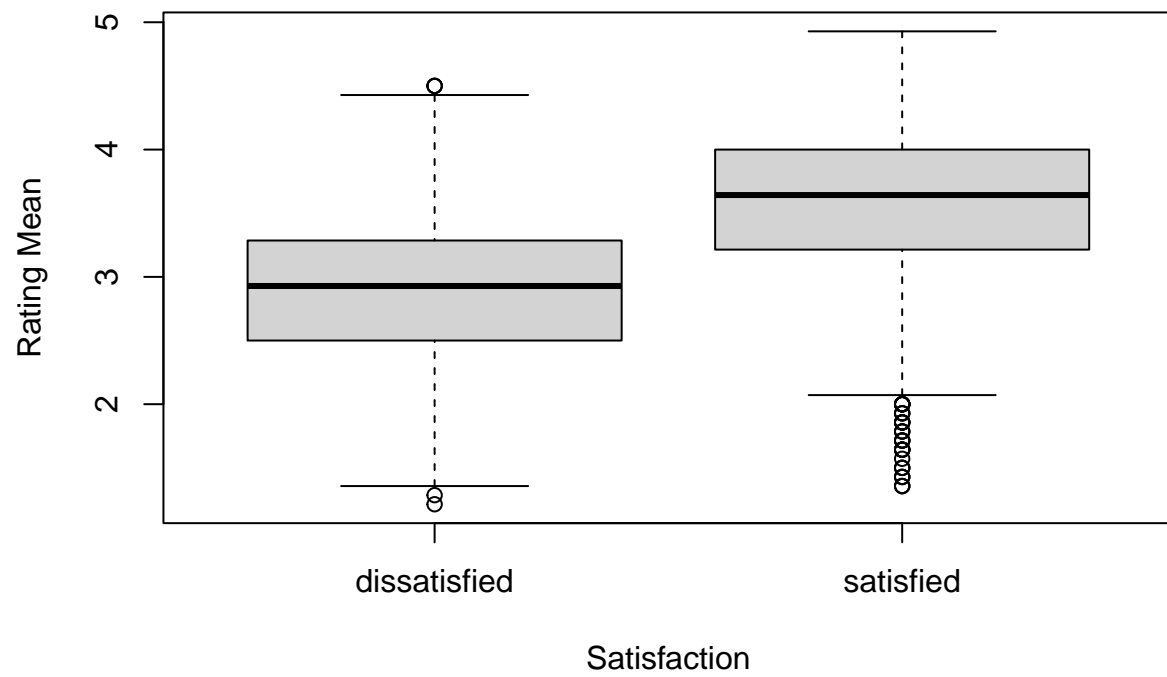
```
dim(train)
```

```
## [1] 6000    3
```


```
summary(train)
```

```
##      satisfaction      Age      rating.mean
## dissatisfied:2775  Min.   : 7.00  Min.    :1.214
## satisfied      :3225  1st Qu.:27.00  1st Qu.:2.786
##                Median :39.00  Median :3.286
##                Mean   :39.14  Mean   :3.294
##                3rd Qu.:51.00  3rd Qu.:3.786
##                Max.   :85.00  Max.   :4.929
```

```
plot(train$satisfaction, train$rating.mean, xlab = "Satisfaction", ylab = "Rating Mean")
```



```
plot(train$satisfaction, train$Age, xlab = "Satisfaction", ylab = "Age")
```




Logistic regression for some baseline

```
glm <- glm(satisfaction~., data=train, family=binomial)
summary(glm)
```

```
##
## Call:
## glm(formula = satisfaction ~ ., family = binomial, data = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.3644  -0.8993   0.3998   0.8531   2.7814
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -6.764123   0.202868 -33.343  < 2e-16 ***
## Age          0.013237   0.001985   6.667 2.61e-11 ***
## rating.mean  1.954379   0.056623  34.515  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 8284.0  on 5999  degrees of freedom
```

```
par(mfrow=c(2,2))
plot(glm)
```



Linear Kernel SVM classification

```
##  
## Call:
```

```
## svm(formula = satisfaction ~ ., data = train, kernel = "linear",
##      cost = 10, scale = TRUE)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##      cost:   10
##
## Number of Support Vectors: 3722
##
## ( 1862 1860 )
##
##
## Number of Classes: 2
##
## Levels:
##   dissatisfied satisfied
```

```
pred <- predict(svm_linear, newdata=test)
table(pred, test$satisfaction)
```


```
##
## pred      dissatisfied satisfied
## dissatisfied      598      248
## satisfied         273      881
```

```
mean(pred==test$satisfaction)
```

```
## [1] 0.7395
```

```
plot(svm_linear, test, rating.mean ~ Age)
```

SVM classification plot



Tune linear kernel SVM

```
tune_svm_linear <- tune(svm, satisfaction~., data=vald, kernel="linear", ranges=list(cost=c(0.001, 0.01
summary(tune_svm_linear)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   0.01
##
## - best performance: 0.2545
##
## - Detailed performance results:
##   cost error dispersion
## 1 1e-03 0.2730 0.03056868
## 2 1e-02 0.2545 0.02763351
## 3 1e-01 0.2565 0.02906411
## 4 1e+00 0.2545 0.02910231
## 5 5e+00 0.2550 0.02915476
## 6 1e+01 0.2550 0.02905933
```

```
## 7 1e+02 0.2555 0.02910231
```

Evaluate linear kernel SVM with best model


```
pred <- predict(tune_svm_linear$best.model, newdata=test)
table(pred, test$satisfaction)
```

```
##
## pred          dissatisfied satisfied
## dissatisfied          564         225
## satisfied            307         904
```

```
mean(pred==test$satisfaction)
```

```
## [1] 0.734
```

```
plot(tune_svm_linear$best.model, test, rating.mean ~ Age)
```



Polynomial Kernel SVM Regression


```
svm_polynomial <- svm(satisfaction~., data=train, kernel="polynomial", cost=10, scale=TRUE)
summary(svm_polynomial)
```

```
##
## Call:
## svm(formula = satisfaction ~ ., data = train, kernel = "polynomial",
##      cost = 10, scale = TRUE)
##
##
## Parameters:
##      SVM-Type:  C-classification
##      SVM-Kernel: polynomial
##          cost:  10
##       degree:  3
##      coef.0:   0
##
## Number of Support Vectors:  4388
##
## ( 2195 2193 )
##
##
## Number of Classes:  2
##
## Levels:
## dissatisfied satisfied
```

```
pred <- predict(svm_polynomial, newdata=test)
table(pred, test$satisfaction)
```


```
##
## pred      dissatisfied satisfied
## dissatisfied      399      117
## satisfied        472     1012
```

```
mean(pred==test$satisfaction)
```

```
## [1] 0.7055
```

```
plot(svm_polynomial, test, rating.mean ~ Age)
```

SVM classification plot



Tune polynomial kernel SVM

```
tune_svm_poly <- tune(svm, satisfaction~., data=vald, kernel="polynomial", ranges=list(cost=c(0.001, 0.1)))
summary(tune_svm_poly)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   0.1
##
## - best performance: 0.302
##
## - Detailed performance results:
##   cost error dispersion
## 1 1e-03 0.3975 0.04124790
## 2 1e-02 0.3260 0.04081122
## 3 1e-01 0.3020 0.03537105
## 4 1e+00 0.3020 0.03417276
## 5 5e+00 0.3020 0.03417276
## 6 1e+01 0.3020 0.03417276
```

```
## 7 1e+02 0.3020 0.03417276
```

Evaluate polynomial kernel SVM

```
pred <- predict(tune_svm_poly$best.model, newdata=test)
table(pred, test$satisfaction)
```


```
##
## pred      dissatisfied satisfied
## dissatisfied      350       94
## satisfied        521      1035
```

```
mean(pred==test$satisfaction)
```

```
## [1] 0.6925
```

```
plot(tune_svm_poly$best.model, test, rating.mean ~ Age)
```

SVM classification plot



Radial Kernel SVM Regression

```
svm_radial <- svm(satisfaction~., data=train, kernel="radial", cost=10, scale=TRUE)
summary(svm_radial)
```

```
##
## Call:
## svm(formula = satisfaction ~ ., data = train, kernel = "radial",
##      cost = 10, scale = TRUE)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##      cost:   10
##
## Number of Support Vectors: 3331
##
## ( 1666 1665 )
##
##
## Number of Classes: 2
##
## Levels:
## dissatisfied satisfied
```

```
pred <- predict(svm_radial, newdata=test)
table(pred, test$satisfaction)
```


```
##
## pred      dissatisfied satisfied
## dissatisfied      611      256
## satisfied         260      873
```

```
mean(pred==test$satisfaction)
```

```
## [1] 0.742
```

```
plot(svm_radial, test, rating.mean ~ Age)
```

SVM classification plot



Tune radial kernel svm

```
tune_svm_radial <- tune(svm, satisfaction~., data=vald, kernel="radial", ranges=list(cost=c(0.1, 1, 10),
summary(tune_svm_radial)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost gamma
##   0.1   0.5
##
## - best performance: 0.2585
##
## - Detailed performance results:
##   cost gamma error dispersion
## 1  1e-01   0.5 0.2585 0.03786306
## 2  1e+00   0.5 0.2605 0.03554731
## 3  1e+01   0.5 0.2625 0.03553168
## 4  1e+02   0.5 0.2630 0.04049691
## 5  1e+03   0.5 0.2620 0.03787406
## 6  1e-01   1.0 0.2590 0.03642344
```

```
## 7  1e+00  1.0 0.2610 0.03580813
## 8  1e+01  1.0 0.2615 0.03793635
## 9  1e+02  1.0 0.2675 0.03795099
## 10 1e+03  1.0 0.2675 0.03466106
## 11 1e-01  2.0 0.2615 0.03682164
## 12 1e+00  2.0 0.2630 0.03902990
## 13 1e+01  2.0 0.2710 0.03921451
## 14 1e+02  2.0 0.2720 0.04029061
## 15 1e+03  2.0 0.2770 0.04679744
## 16 1e-01  3.0 0.2620 0.03780065
## 17 1e+00  3.0 0.2690 0.03885872
## 18 1e+01  3.0 0.2730 0.04070217
## 19 1e+02  3.0 0.2715 0.04743709
## 20 1e+03  3.0 0.2785 0.05452064
## 21 1e-01  4.0 0.2640 0.03820995
## 22 1e+00  4.0 0.2690 0.04005552
## 23 1e+01  4.0 0.2705 0.04615493
## 24 1e+02  4.0 0.2775 0.04837642
## 25 1e+03  4.0 0.2925 0.04739022
```

Evaluate radial kernel svm with best model


```
pred <- predict(tune_svm_radial$best.model, newdata=test)
table(pred, test$satisfaction)
```

```
##
## pred      dissatisfied satisfied
## dissatisfied      583      242
## satisfied        288      887
```

```
mean(pred==test$satisfaction)
```

```
## [1] 0.735
```

```
plot(tune_svm_radial$best.model, test, rating.mean ~ Age)
```



How Results Were Achieved?

For my classification data, radial kernel SVM worked the best marginally compared to linear and polynomial. This could be due to the nature of the data not being able to be separated linearly in a clean manner. However, it should be noted that the difference between my radial and linear results were very minimal. This could be the direct result of me aggregating the features that have a rating from 1 to 5 into one called the rating mean. In addition, the tuning process did not turn out as expected, again possibly due to the aggregation step that was done at the beginning.